

---

# Matrix Library Implementation

Juliet James & Trisha Kholiya

---

## Our Goals

- Build a lightweight C++ matrix library
- Support core matrix operations
- Implement an eigensolver for symmetric matrices
- Validate accuracy + compare performance with Armadillo
- Use our library as a drop-in replacement for HW3

# Library Architecture

---

## Matrix Class

- Storage: flat vector
- Operators: + - \* scalar multiply
- Utilities: transpose, identity, zeros, ones, random
- Eigensolver Pipeline

Householder → Tridiagonal T



QL → Eigenvalues



Combine  $Q_{\text{house}} \times Q_{\text{ql}}$  → eigenvectors



# Constructors

```
Matrix(int rows, int cols);  
  
Matrix();  
  
Matrix(const vec& values, int rows, int cols); LVALUE  
  
Matrix(vec&& values, int rows, int cols); RVALUE  
  
static Matrix Ones(int rows, int cols);  
  
static Matrix Zeros(int rows, int cols);  
  
static Matrix Random(int rows, int cols);  
  
static Matrix Identity(int n);
```



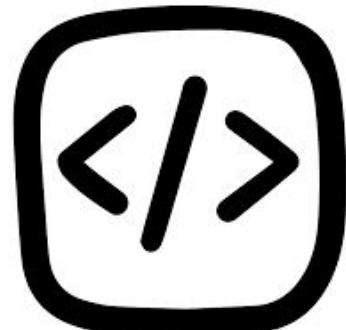
# Operators

```
double& operator()(int x, int y);  
const double& operator()(int x, int y) const;  
bool operator==(const Matrix& other) const;  
Matrix operator+(const Matrix& other) const;  
Matrix operator-(const Matrix& other) const;  
Matrix operator*(const Matrix& other) const;  
Matrix operator*(double s) const;  
std::ostream& operator<<(std::ostream & out, const Matrix & M);
```

---

## Optimizing Performance in our Library

- STORAGE: Switched to 1D storage
- Optimized our operators by using raw pointers as much as possible -> direct memory access
- unit-stride vector operations (i, i +1, etc.) optimizes performance



# Multiplication Code

---

```
for (int i = 0; i < rows; ++i) {  
    for (int j = 0; j < cols; ++j) {  
        double sum = 0.0;  
        for (int k = 0; k < my_cols; ++k) {  
            sum += m1[i * my_cols + k] * m2[k * cols + j];  
        }  
        r[i * cols + j] = sum;  
    }  
}  
  
for (int i = 0; i < rows; ++i) {  
    for (int k = 0; k < my_cols; ++k) {  
        double a = m1[i * my_cols + k];  
        for (int j = 0; j < my_cols; ++j) {  
            r[i * cols + j] += a * m2[k * cols + j];  
        }  
    }  
}
```

- Removed temporary variable sum
- Decreased direct memory access
- Increased performance!

# Householder & QL Overview

---

## Householder tridiagonalization

- Uses reflections to zero out lower entries
- Produces tridiagonal  $T$
- Stores transforms in  $Q_{\text{house}}$
- Makes eigen-solve cheap ( $O(n^2)$ )

## QL on $T$

- Implicit-shift iterations
- diagonalize  $T$
- Diagonal  $\rightarrow$  eigenvalues
- Stores transforms in  $Q_{\text{ql}}$

$$\begin{bmatrix} \times & \times & \times \\ \times & \times & \times \end{bmatrix} \xrightarrow{Q_1} \begin{bmatrix} \times & \times & \times \\ 0 & \times & \times \end{bmatrix} \xrightarrow{Q_2} \begin{bmatrix} \times & \times & \times \\ \times & \times & \times \\ 0 & \times & \times \\ 0 & \times & \times \\ 0 & \times & \times \end{bmatrix} \xrightarrow{Q_3} \begin{bmatrix} \times & \times & \times \\ \times & \times & \times \\ \times & \times & \times \\ 0 & \times & \times \\ 0 & \times & \times \end{bmatrix} \quad (10.1)$$

$A$                      $Q_1 A$                      $Q_2 Q_1 A$                      $Q_3 Q_2 Q_1 A$

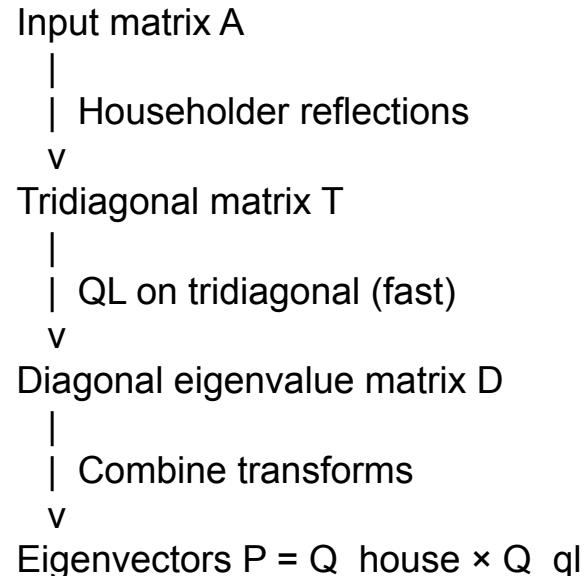
Each reflection wipes out one column of lower entries until only a tridiagonal band remains.

Trefethen, L. N., & Bau, D. (1997).  
Numerical Linear Algebra. Philadelphia:  
SIAM.

# How the Pieces Combine ( $A \rightarrow T \rightarrow D$ )

---

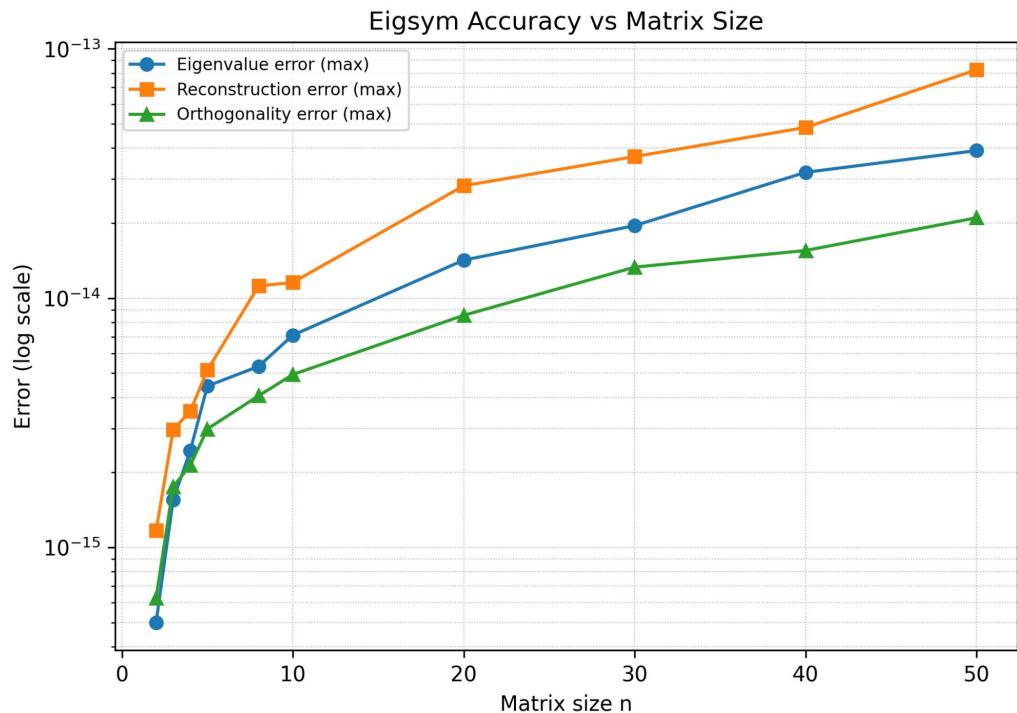
- Householder:  $A \rightarrow T$
- QL:  $T \rightarrow$  eigenvalues + local eigenvectors
- Full eigenvectors:  $P = Q_{\text{house}} \times Q_{\text{ql}}$
- Validation:  $A \approx P \Lambda P^T$



# Accuracy Validation & Results

---

- Reconstruction:  $A \approx P \Lambda P^T$
- Orthogonality:  $P^T P \approx I$
- Perfect on identity + diagonal tests
- Errors grow with matrix size (floating-point + iterative QL)
- Stable and accurate for small–medium matrices



---

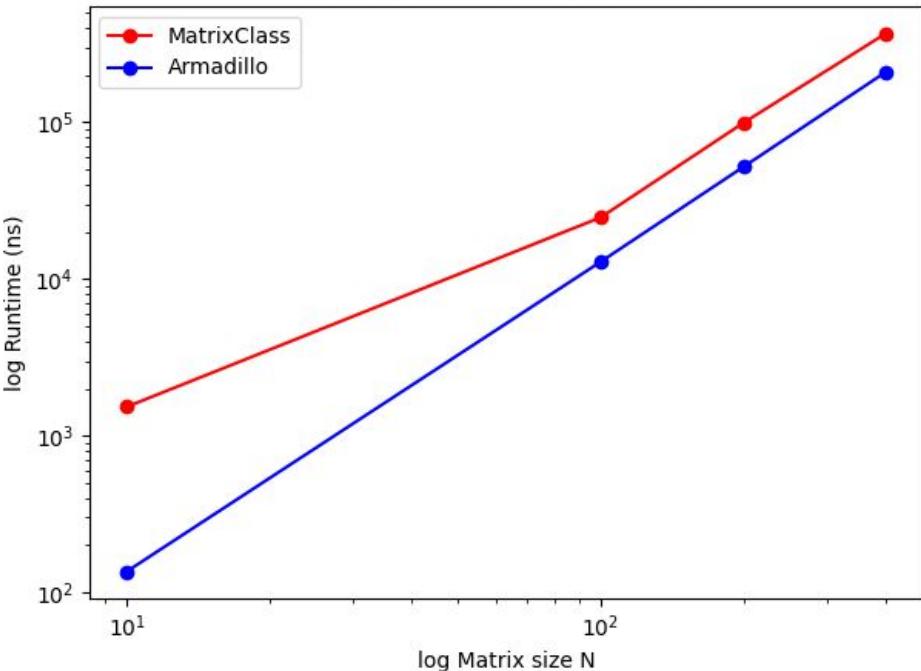
## Accuracy vs Armadillo (why ours diverges)

- Armadillo uses highly optimized LAPACK routines
- Better numerical stability + shift strategies
- Our solver works but isn't as optimized
- Produces correct eigenpairs for moderate sizes

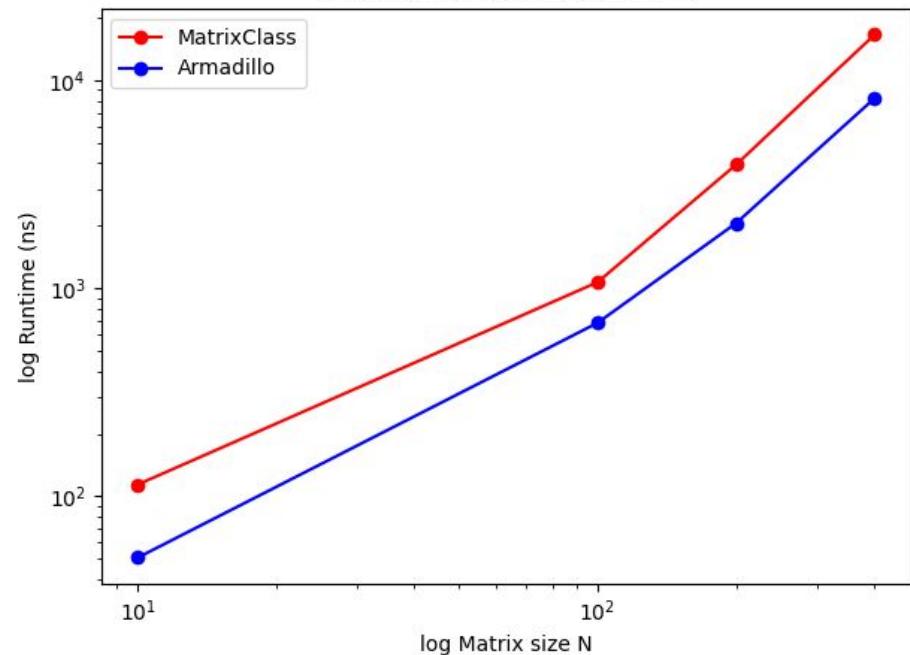
# Benchmarking: Constructors

Google benchmark

InitializationRandom Benchmark



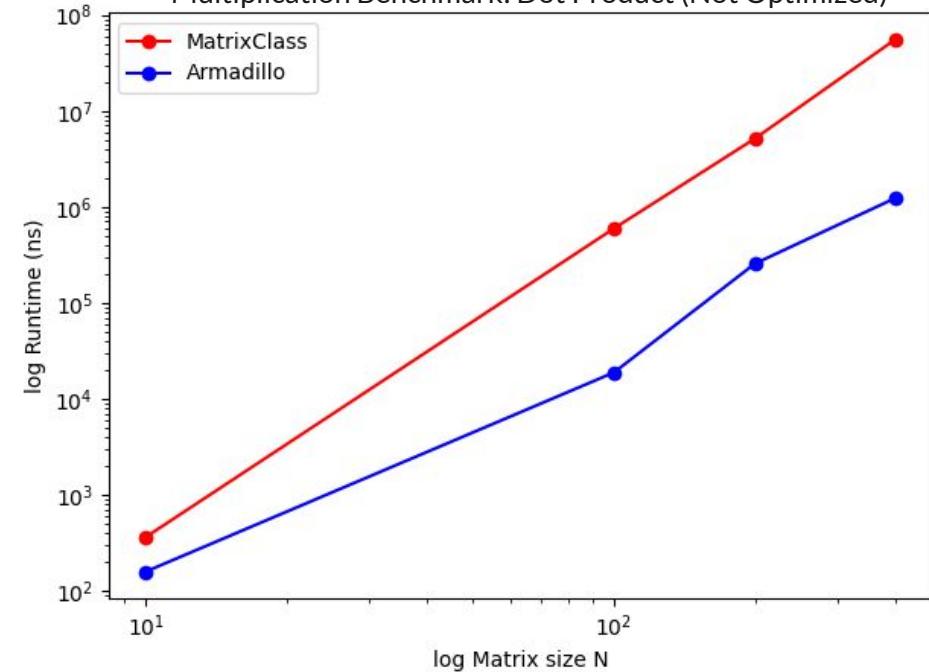
InitializationZeros Benchmark



# Benchmarking: Multiplication

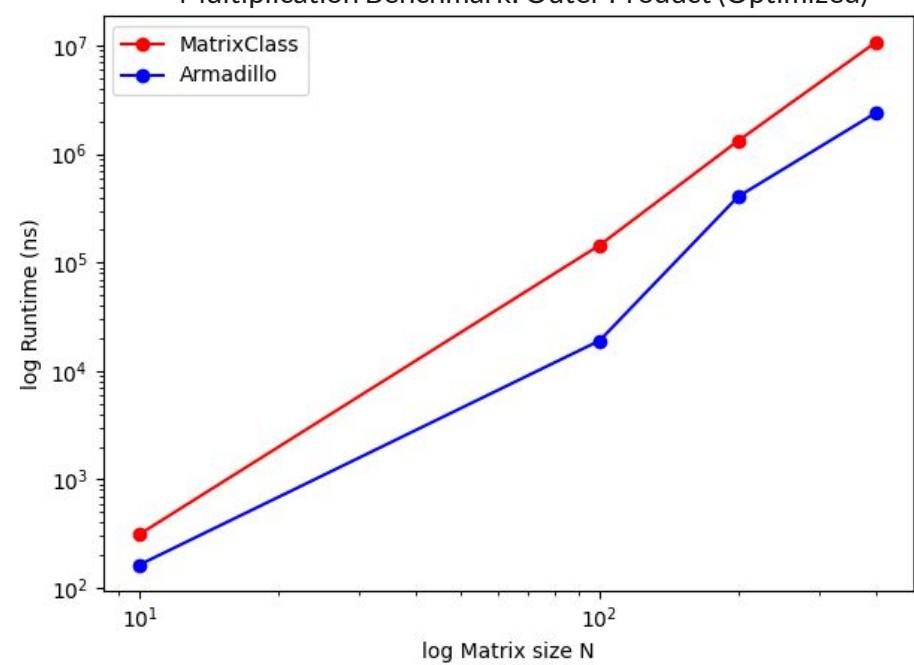
BEFORE

Multiplication Benchmark: Dot Product (Not Optimized)



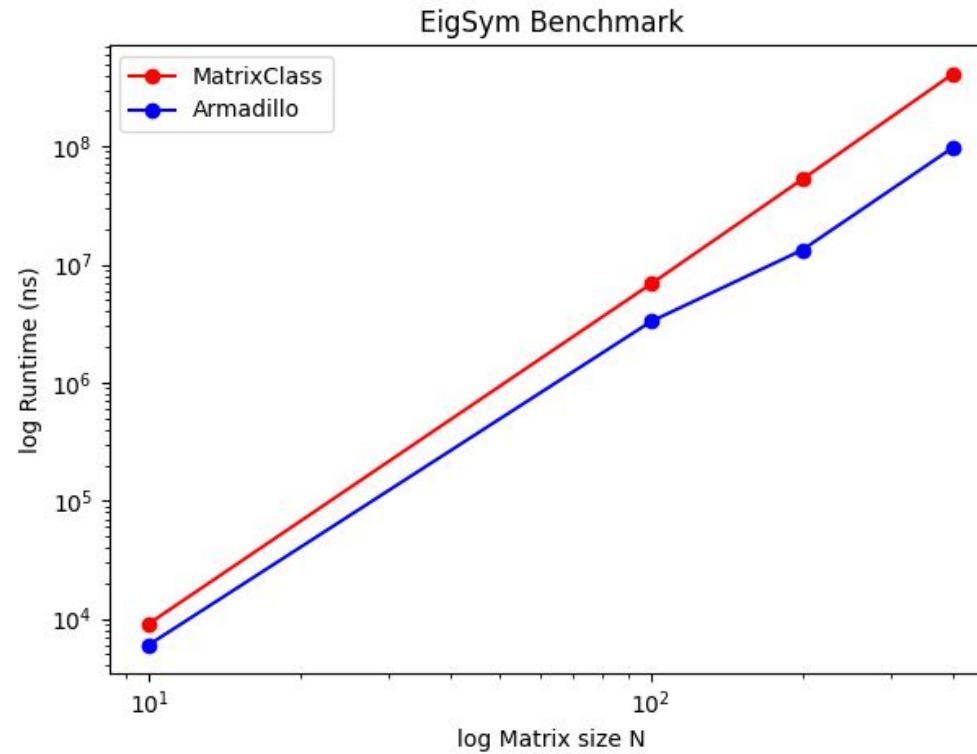
NOW

Multiplication Benchmark: Outer Product (Optimized)



# Benchmarking EigSym

---



---

## Armadillo is Faster..why?

- Armadillo uses high-performance SIMD ->
- Can be used with NVBLAS to obtain GPU-accelerated matrix multiplication
- Delayed Evaluation: Expression Templates
  - Lightweight marker objects that hold references to matrices and data associated with specific operations.
  - The marker objects can be chained together, leading to the full description of an arbitrary -> reduces temporary variables

---

# Packaging the Library

- CMakeLists includes Install Section and Lib includes a config file so that the package can be installed by users that need Matrix Class
- Custom Exception class to make errors more user readable
- Printing vector class

---



## Replacing Armadillo in HW3

DEMO!

---

## Next Steps

- Implement HPC methods to boost performance
- Explore more optimization opportunities i.e. delayed evaluation
- Extend Matrix class to have more vector-capabilities like Armadillo
- Generalization to non-symmetric matrices is a key next milestone



# References

Conrad Sanderson and Ryan Curtin.

Armadillo: An Efficient Framework for Numerical Linear Algebra.  
International Conference on Computer and Automation Engineering, 2025.  
<https://doi.org/10.1109/ICCAE64891.2025.10980539>

Conrad Sanderson and Ryan Curtin.

Practical Sparse Matrices in C++ with Hybrid Storage and Template-Based Expression Optimisation.  
Mathematical and Computational Applications, Vol. 24, No. 3, 2019.  
<https://doi.org/10.3390/mca24030070>

Trefethen, L. N., & Bau, D. (1997).

Numerical Linear Algebra. Philadelphia: SIAM.

<https://pubs.siam.org/doi/book/10.1137/1.9781611977165>

John D. McCalpin

A Survey of Memory Bandwidth and Machine Balance in Current High Performance Computers

<https://www.cs.virginia.edu/~mccalpin/papers/balance/>

Daniel Lemire

Rolling your own fast matrix multiplication: loop order and vectorization

[https://lemire.me/blog/2024/06/13/rolling-your-own-fast-matrix-multiplication-loop-order-and-vectorization/?utm\\_source=chatgpt.com](https://lemire.me/blog/2024/06/13/rolling-your-own-fast-matrix-multiplication-loop-order-and-vectorization/?utm_source=chatgpt.com)

---



# Thank you :)

Please use our Matrix Library for your code!!!